

ING Credit Infrastructure: working simulation

Product Owner & Senior Risk Data Journey Expert · REQ-10118811

Prep for Anjish Bhondwe. Ecosystem map (tribes, chapters, squads, interactions) reconstructed from the public JD, sibling Credit Infrastructure postings, and public ING Belgium Credit Risk signals. Not an internal ING org chart. Labels: **JD** = named in posting · **Public** = public LinkedIn / careers · **Sim** = realistic Orange House fill-in.

1. Context snapshot

Where: ING Belgium, Avenue Marnix · Integrated Risk Management · Risk Platforms · Credit Infrastructure. JD

Domain: Retail and business banking credit risk (not Investment Banking). Engines feed capital, RWA and LLP.

Engines in landscape: rule-based (DoD, EWS, and related: DSTI, DTAI, LTV) and model-driven (PD, EAD, LGD). Plus model data mart. JD / sibling CR Data Engineer

Your paper role: 50% Product Owner, 50% Senior Risk Data Journey Expert. Bridge between Credit Risk business and operational tribes (especially Data Management), plus tooling for the Risk family and Business tribes. JD

2. Where you sit (department stack)

CRO / Risk Belgium → **Integrated Risk Management (IRM)** → **Risk Platforms** (domain) → **Credit Infrastructure** (your team / squad home) → Agile squad delivering credit engines + risk data journey work.

- **IRM:** connects risk methodologies, platforms and oversight so Belgium can run regulated credit risk at scale.
- **Risk Platforms:** builds and runs the systems under risk (engines, data paths, tooling), not credit policy itself.
- **Credit Infrastructure:** batch engines + model data mart + lake regulatory consumption for credit risk parameters.
- **You:** own squad backlog and master the credit risk data journey so engines stay policy-aligned and auditable.

3. Quick map (detail on next pages)

Inside your squad: you (PO + Risk Data Journey Expert) · Credit Risk Data Engineers (BE + PL) · functional analysts / modeler partners · often a Scrum Master / Agile Coach shared. **Outside:** Credit Risk Tribe · Data Management · Business tribes · Finance · Audit · chapter leads. Full squad functions and the credit risk data journey landscape follow.

3a. Your Credit Infrastructure squad: people and functions

Simulation of a typical Credit Infrastructure agile squad. Headcount and titles vary; functions are what interviewers expect you to orchestrate. JD + sibling CR Data Engineer + Orange House sim

Person / seat	Core function in the squad	Typical week	Hands off to you when...
You: Product Owner (50%)	Own and prioritise the product backlog; turn Credit Risk / regulatory priorities into delivery plans; protect capacity; escalate cross-tribe blockers; accept or reject done work against criteria	Refinement, planning, stakeholder syncs, SteerCo status, trade-off calls	Always on: priorities and “what ships this sprint”
You: Senior Risk Data Journey Expert (50%)	Master the Credit Risk Data Landscape (PD, EAD, LGD + related engines); design end-to-end data journeys; lake regulatory sequencing; flows/tooling for Risk and Business; digitalization; auditor key contact	Lineage walkthroughs, lake-engine epic design, audit prep, tooling thin slices	Journey design and landscape truth (source → lake → engine → capital)
Credit Risk Data Engineer (Brussels)	Build/maintain batch engines and jobs (Oracle, SAS EG / DI patterns); implement engine logic; run and stabilise overnight batches; support mart; debug joins and reconciliations	Coding, batch ops, peer review, incident triage, story estimates	Technical feasibility, runtime risk, “data vs engine” diagnosis
Credit Risk Data Engineer (ING Hubs / Warsaw)	Same craft as BE engineers: engines, landscape pipelines, tooling for Risk and Business tribes; extend capacity without hallway context	Remote ceremonies; pick up well-written stories; pair on complex lineage	Needs clear acceptance, samples, and glossary from you
Credit Risk Functional Analyst often embeds or partners	Translate policy and model intent into functional requirements; define expected indicator behaviour; support UAT/acceptance with Credit Risk	Spec workshops, acceptance cases, edge cases (forbearance, DoD flags)	Ambiguous policy language that would block engineers
Modeler / Risk Analytics partner mart consumer	Need model data mart features for development and monitoring; challenge data fitness for PD/EAD/LGD use	Mart requests, monitoring checks, model change impact	Mart backlog pressure vs engine stability trade-offs
Scrum Master / Agile Coach sim · often shared	Facilitate ceremonies; remove process blockers; keep flow visible; coach team health	Stand-ups, retros, impediment board	Does not own backlog priority (you do)

How the squad divides a story

1. **You (journey):** map the hop (which source, lake field, engine, output, consumer).
2. **You (PO):** cut the epic into lake dependency + engine story + acceptance; sequence on the board.
3. **Functional analyst / Risk SME:** expected behaviour and edge cases.
4. **Data engineer(s):** implement, test, run batch, reconcile.
5. **Credit Risk:** accept parameter impact.
6. **You:** close story; stash audit evidence if capital-relevant.

What “done” means for this squad

- Engine or flow change matches accepted behaviour (not only code merged).
- Reconciliation or control check passed on the agreed portfolio slice.
- Dependencies (especially Data Lake) are explicit: done or consciously deferred.
- If capital/LLP relevant: lineage note ready for auditor questions.

3b. Credit Risk Data Landscape (what “landscape” means)

The landscape is the map of credit risk data and engines that produce parameters for capital, RWA and LLP. Your journey-expert job is to know how each piece connects, not to memorise every table on day one.

Landscape layer	What lives here	Why it matters
Sources / operational data	Lending products, facilities, collateral, customer behaviour, arrears, forbearance flags; residual Record Bank codes where still present	Garbage in → wrong DoD/PD/LGD
Data Lake (regulatory initiatives)	Governed landing of attributes required by regulation or policy; quality rules; lineage start	JD: you support regulatory initiatives here with Data Management
Rule-based engines	DoD (definition of default), EWS (early warning), related indicators (DSTI, DTAL, LTV, etc.)	Policy logic turned into repeatable batch calculation
Model-driven engines	PD, EAD, LGD (and related model outputs)	Feed RWA, capital, LLP; model version sensitive
Model data mart	Curated data for developing and monitoring new models	Bridge to modelers / validation; JD names this duty
Outputs & tooling	Parameters and flows used by Risk family and Business tribes; reports; controls	Digitalization and tooling stories live here
Downstream consumers	Capital / RWA / LLP, portfolio steering, underwriting/monitoring, Audit	Your deadlines and evidence pack

3c. Credit risk data journey (end to end)

Risk / regulatory need → Source readiness → Data Lake land + quality → Engine input contract → Rule or model engine run → Reconciliation / controls → Parameter publish → Capital / RWA / LLP + Risk & Business tooling → Auditor evidence			
Hop	What happens	Who leads	Your journey-expert move
1. Need	Policy or regulator requires a field, rule, or model behaviour by date D	Credit Risk Tribe	Capture why, date, portfolio scope, acceptance language
2. Source	Confirm attribute exists in operational systems (or Record residual map)	Data Management + Business ops	Flag gaps early; do not start engine story blind
3. Lake	Land attribute, quality rule, catalogue / lineage	Data Management	Sequence lake story before engine consumption
4. Contract	Define engine inputs, joins (e.g. facility-collateral), defaults, history	You + functional analyst + engineer	Write the journey map for the story
5. Engine	Batch computes DoD/EWS/PD/EAD/LGD (or related)	Data engineers	Protect capacity; clarify rule vs model path
6. Control	Reconcile counts, rates, breaks vs prior run or benchmark	Engineers + Risk	Insist acceptance includes a control check
7. Publish	Parameters available to capital, LLP, tooling, Business indicators	Credit Infrastructure + consumers	Notify Finance/Risk if window risk
8. Evidence	Lineage pack for internal/external audit	You (key contact) + engineer + Risk SME	Keep boring, complete evidence; backlog gaps

3d. Journey by engine family (interview showcase)

Engine / indicator	Journey focus	Typical consumers	Squad function that owns build
DoD	Policy events (arrears, forbearance, unlikely-to-pay) → flags on facilities/customers	Credit Risk, capital/LLP, monitoring	Data engineer implements; FA + Risk accept
EWS	Early signals from behaviour / financial stress → watchlist style outputs	Credit Risk monitoring, Business follow-up	Engineer; Risk tunes thresholds with you
DSTI / DTAI / LTV	Affordability and collateral ratios from income, debt, property values	Business lending journeys, Credit Risk	Engineer; Business tribe feels UX impact
PD	Governed exposures + behaviour → probability of default (model)	RWA, capital, portfolio steering	Engineer + modeler; you own journey fitness
EAD	Drawn/undrawn and conversion factors → exposure at default	RWA, capital	Same; facility data quality critical
LGD	Collateral, recovery, cure paths → loss given default	LLP, RWA	Joins (facility-collateral) often the pain point
Model data mart	Curated history for model build/monitoring	Modelers, validation	Engineers deliver mart features you prioritise

Landscape view you can draw on a whiteboard in 60 seconds

Sources (lending / collateral / behaviour / Record residuals)
→ Data Lake (regulatory attributes + quality) [Data Management]
→ Credit engines: DoD · EWS · DSTI/DTAI/LTV · PD · EAD · LGD [your squad]
→ Model data mart [your squad + modelers]
→ Outputs to capital / RWA / LLP · Risk tooling · Business indicators
→ Audit evidence path [you as key contact]

Squad swimlane on one journey (example: new forbearance field → PD)

Step	You (PO)	You (journey)	Data engineer	FA / Risk	Data Mgmt
Scope	Prioritise epic	Draw hops	Estimate	Policy meaning	Field exists?
Lake	Sequence board	Acceptance for land	Consume later	Confirm definition	Land + quality
Engine	Capacity guard	Input contract	Code + batch	Expected PD impact	Support lineage
Done	Close / escalate	Evidence note	Reconcile	Accept	Catalogue update

This is the "credit risk data journey under the landscape": landscape = the map of layers and engines; journey = how one need travels those layers with clear owners.

4. Tribes: who they are and what they must be doing

ING's Orange House model: tribes own a value domain; chapters own craft across squads. Names below mix JD anchors with public Belgium Credit Risk signals and typical adjacent tribes you will meet every week.

A. Credit Risk Tribe Public · Gregory Dekimpe, Credit Risk Tribe lead ING Belgium

Mission (public language): end-to-end credit risk transformation: policies and decision models, analytical insights for portfolios, future-proof data/platforms for sustainable growth.

What they must be doing:

- Set and maintain credit policies and risk appetite for retail and business banking books.
- Own definition of default, forbearance, early warning logic at policy level (you implement engines).
- Steer portfolios with data (portfolio steering / PI-style work sits here or adjacent under the tribe).
- Accept engine outputs that affect underwriting, monitoring, LLP narrative, and capital conversations.
- Area leads by segment (e.g. Business Banking credit risk area lead) translate tribe strategy into segment practice. Public

How they interact with you: They are your primary “why.” Policy change → backlog item. Portfolio noise → engine or data defect triage. They sit in refinement and acceptance; you do not rewrite policy for them.

B. Your home: Risk Platforms / Credit Infrastructure (under IRM) JD

What the tribe/domain must be doing: keep risk calculation and risk-data platforms reliable, compliant, and moving toward target architecture (Data Lake + digital tooling), while legacy Oracle/SAS batches still run.

Your squad's slice: credit engines (DoD, EWS, PD, EAD, LGD, related indicators), model data mart, regulatory lake consumption, flows/tools for Risk and Business.

How other Risk Platforms squads interact (sim): shared platform standards, release calendars, lake patterns, maybe market/ops risk platforms nearby. You sync on dependencies, not on credit policy.

C. Data Management tribe (operational) JD names this tribe

What they must be doing: own operational data products and the Data Lake: landing zones, quality rules, catalogues, access, and often the pipelines that make regulatory fields available bank-wide.

How they interact with you: almost every regulatory initiative in the JD hits them first. You sequence: lake readiness → engine consumption. Joint refinement when a field is missing or joins break (facility-collateral, customer-exposure).

Failure mode: you promise Risk a PD/DoD change before Data Management can land the attribute. Fix: make the dependency visible on the backlog.

D. Business tribes (Retail / Private Individuals, Business Banking, and related) JD: “Business tribes”

What they must be doing: run customer and product value streams (accounts, lending origination, servicing). They care that risk engines and indicators (LTV, DSTI, EWS, DoD flags) do not block healthy growth or create unfair frictions.

How they interact with you: less daily than Credit Risk, but they consume tooling and indicators. Digitalization stories (replace spreadsheet controls with flows) often start as Business or Risk pain and land on your backlog as thin tooling slices.

E. Finance / CFO reporting adjacency Sim · capital / RWA / LLP consumers

What they must be doing: capital, RWA, provisions and regulatory reporting calendars. They are downstream of your parameters.

How they interact with you: reporting windows become your hard deadlines. When engines slip, Finance feels it. You escalate early with impact language (which parameter, which report, which date).

F. Model Risk / Model Validation / Model Strategy adjacency Sim · group & local model life cycle

What they must be doing: challenge and oversee IRB / IFRS9 model use, monitoring methodologies, and model life cycle governance. Local modelers and your mart feed that world.

How they interact with you: mart features for monitoring/validation; evidence when an engine parameter is model-driven; change control when model versions bump. You do not own model approval. You own that data and engines support approved use.

G. Audit & Compliance (2nd/3rd line) JD: auditor key contact

What they must be doing: test controls, lineage, and regulatory compliance around capital-relevant engines.

How they interact with you: you are the named bridge. Prepare journey maps, reconciliations, and walkthroughs. Backlog control gaps. Do not improvise policy answers: bring Credit Risk for “why,” engineers for “how.”

H. ING Hubs (Poland) delivery capacity Sibling CR Data Engineer Warsaw

What they must be doing: supply Credit Risk Data Engineers into the same Credit Infrastructure work (engines, landscape, tooling for Risk and Business tribes).

How they interact with you: remote squad members in ceremonies; clear stories and acceptance; timezone-aware refinement. Your PO craft shows up in how cleanly Warsaw can pick up work without tribal knowledge only in Brussels.

5. Chapters: craft homes (across tribes)

Chapters cut across squads. People report for craft quality here; they deliver in squads. Names are Orange House style sims aligned to JD roles.

Chapter	Who sits there	What the chapter must be doing	Touch with you
Product / Customer (Risk Data) Journey Expert chapter	POs, CJses, journey experts in risk/data	Backlog quality, journey mapping standards, stakeholder management craft, hiring bar	Your chapter home; coaching, peer review of journeys
Data Engineering chapter	Credit Risk Data Engineers, lake/pipeline engineers	Coding standards, Oracle/SAS/DI patterns, lake pipelines, DoD for data stories	Engineers’ craft home; you align story DoD with them
Risk Analytics / Modelling chapter	Modelers, functional credit risk analysts	Model monitoring craft, functional specs, analytics methods	Mart and engine functional clarity
Credit Risk Management chapters (in Credit Risk Tribe)	Policy, segment CRM, analytics, restructuring (pattern known publicly in ING CRM tribes)	Policy craft, portfolio analytics, segment expertise (PI, BB, PB/WM patterns)	Demand quality and acceptance; not your people-manager line
Data Management chapters	Data stewards, lake POs, quality engineers	Data product ownership craft, quality SLAs, catalogue discipline	Joint standards on lineage and field definitions

Chapter vs squad rule for interviews: squad owns delivery and deadlines. Chapter owns how craft is done. Never put a capital deadline into “chapter theatre.”

6. Squads near Credit Infrastructure (simulation map)

Squad (sim name)	Likely focus	Interaction with your squad
Credit Engines / Credit Infrastructure (yours)	DoD, EWS, PD/EAD/LGD and related batch engines; mart; lake regulatory consumption	You lead as PO
Model Data Mart / Model Support	Mart features for model development and monitoring	Shared backlog pressure; often same engineers or sibling squad
Data Lake Regulatory / Credit Risk Data Products	Landing regulatory attributes for risk use	Hard dependency; joint epics
Credit Risk Reporting	Risk reports and dashboards on portfolio / engine outputs	Hand off outputs; avoid owning their dashboards
Portfolio Credit Risk / Portfolio Steering	Insights and steering under Credit Risk Tribe Public tribe signal	Consume parameters; raise data quality issues back to you
Lending / Credit Decisioning (Business tribe)	Origination and decision journeys using LTV/DSTI/EWS signals	Indicator contracts; change impact assessment

7. How tribes interact with you (operating model)

Demand path

Credit Risk Tribe (policy / portfolio need) → you (clarify + backlog) → Data Management (if lake/source) → Credit Infrastructure engineers (build) → Credit Risk acceptance → Finance/Audit consume evidence

Who brings what into refinement

- **Credit Risk:** policy text, effective date, portfolio impact, acceptance criteria in risk language.
- **Data Management:** field availability, quality, lineage, capacity to land attributes.
- **Engineers (BE/PL):** technical slice, engine runtime risk, regression scope.
- **You:** one backlog item with clear owner of each dependency; cut scope if capacity is false.
- **Modelers:** mart or monitoring need; must not silently become ten “P0” features.

Standing forums (realistic week)

Forum	Who	Your job
Squad ceremonies	Engineers, SM/facilitator, you	Priorities, blockers, capacity honesty
Risk / Data / IT sync	Credit Risk, Data Management, IT, you	Surface cross-tribe dependencies early
Chapter huddle	PO/CJE peers + chapter lead	Craft only; bring patterns, not panic deadlines
Incident triage	Engineers + Risk on-call style	Impact language (LLP/RWA/report date)
Audit prep / walkthrough	Audit, Risk, engineer, you	Key contact; journey map + evidence pack
SteerCo / tribe update	Tribe/domain leads	Two-sentence status + risk + ask

RACI-style (interview-ready)

Topic	Credit Risk Tribe	You (CI)	Data Mgmt	Engineers	Audit
Policy / DoD definition	A/R	C	I	I	I
Engine backlog priority	C	A/R	C	C	I
Lake field landing	C	C	A/R	C	I
Engine build / batch run	I	C	I	A/R	I
Acceptance of parameter	A/R	C	I	C	I
Lineage evidence pack	C	A/R	C	R	A (review)

A = accountable · R = responsible · C = consulted · I = informed. Soft model for interviews, not an official ING RACI.

8. Interaction examples (what “good” looks like)

- **Credit Risk → you:** “Forbearance attribute X must drive DoD/PD from date D.” You translate into lake story + engine story with one acceptance chain.
- **You → Data Management:** “We need attribute X in lake with quality rule Q by sprint N-1, or engine story slips.”
- **You → Warsaw engineer:** story with inputs, expected output sample, reconciliation check, rollback note.
- **Portfolio steering → you:** “Segment Y LLP noise.” You open lineage: source vs engine vs report path.
- **Business tribe → you:** “EWS floods branch noise.” You check indicator threshold vs data defect vs policy intent with Credit Risk.
- **Audit → you:** “Show RWA-driving parameter path.” You run the walkthrough with engineer + Risk SME.

9. Data journey pointer

Full Credit Risk Data Landscape and 8-hop journey are in sections **3b to 3d** (layers, hops, engine families, squad swimlane). One-line reminder: Source → Lake → Engines (DoD/EWS/PD/EAD/LGD + related) → Mart / tooling → Capital · RWA · LLP → Audit evidence.

10. Problem scenarios (cross-tribe)

Scenario A: Regulatory field must land in the Data Lake

Tribes: Credit Risk (need) · Data Management (land) · Credit Infrastructure (consume) · Finance (window).

Your move: one epic with two stories (lake then engine); protect sequence; escalate if lake slips before Risk deadline.

Scenario B: Broken facility-collateral join (LGD / LLP)

Tribes: Data Management + your engineers + Credit Risk + Finance impact.

Your move: separate data defect vs engine bug; assign owner; protect sprint capacity; document lineage for audit.

Scenario C: Backlog conflict (policy vs tech debt vs mart)

Tribes: Credit Risk · Risk Analytics/modelers · your chapter craft vs delivery reality.

Your move: mandatory regulatory first; one mart story that unblocks; park the rest visibly.

Scenario D: Auditor asks for RWA-driving parameter evidence

Tribes/lines: Audit · Credit Risk · engineers · you as key contact.

Your move: journey map walkthrough; backlog control gaps; no blame theatre.

Scenario E: Record Bank residual mapping

Tribes: Credit Risk + Data Management historical products · your landscape role.

Your move: map residual codes with SMEs; backlog fix; honest ramp (no fake Record mastery).

Scenario F: Business tribe tooling ask vs Risk digitalization

Tribes: Business tribe pain · Credit Risk consumption · your digitalization duty.

Your move: thin automated flow on engine outputs; clear acceptance; do not build a full Business reporting product owned by another squad.

Scenario G: Warsaw capacity vs Brussels tribal knowledge

Tribes: Hubs delivery · local Credit Risk context.

Your move: write stories hubs can execute; pair on first lake+engine epic; keep glossary of Record/ING codes in the squad wiki.

11. Challenges you should expect

- Legacy Oracle/SAS batch plus Data Lake target architecture at the same time.
- Regulatory calendars that ignore sprint boundaries.
- Many tribes, one squad capacity (Risk, Data, Business, Finance, Audit).
- Auditor evidence must be boringly complete.
- 50/50 split: deep landscape work vs ceremony and backlog ownership.
- Chapter craft pressure vs tribe delivery pressure: protect both without mixing them.

12. First 90 days (on the job)

1. Map tribes: who is Credit Risk Area SME, Data Management PO, chapter leads, Warsaw leads.
2. Map PD / EAD / LGD and related engine flows end to end with engineers and Risk.
3. Deliver or materially advance one Data Lake regulatory initiative with Data Management.
4. Land one flow/tooling improvement for Risk or Business.
5. Earn squad trust: clear priorities, honest capacity, early escalation.
6. Be ready as auditor contact on one engine/control topic.

13. Hiring manager Q&A (domain + org)

Q1. What do you think Credit Infrastructure does?

A. It develops and maintains the engines that calculate credit-risk indicators and parameters that feed capital, RWA and LLP. It keeps data infrastructure for future risk models and bridges Credit Risk, Data Management and operational teams. That matches the JD.

Q2. Name the tribes you expect to work with.

A. Credit Risk Tribe as demand and acceptance. Data Management as the operational tribe for lake and quality. Business tribes as consumers of indicators and tooling. My home is Risk Platforms / Credit Infrastructure under IRM. Finance and Audit sit adjacent on capital and evidence.

Q3. How do chapters differ from tribes here?

A. Tribes own value and priorities in a domain. Chapters own craft across squads (PO/journey, data engineering, risk analytics). I deliver in the squad for the tribe/domain; I improve craft in the chapter.

Q4. How do you see the 50/50 split in a normal week?

A. Half on backlog, ceremonies, dependencies and stakeholder alignment across tribes. Half on landscape: PD/EAD/LGD journeys, lake regulatory work, flows and tooling. When a regulatory window hits, both halves collapse into one priority list.

Q5. Walk me through PD, EAD and LGD in this landscape.

A. They sit in the model-driven layer of the Credit Risk Data Landscape. Journey: governed exposures and behaviour land via lake/sources, engines compute PD/EAD/LGD, controls reconcile, then parameters feed RWA and LLP. I have worked that chain at HSBC, Credit Suisse and on calculation engines at KBC.

Q5b. What is the difference between landscape and journey?

A. Landscape is the map of layers and engines (sources, lake, rule engines, model engines, mart, outputs). Journey is how one regulatory or business need travels those layers with clear owners and acceptance. My JD is to master both.

Q5c. Who does what inside your squad on a lake-to-engine change?

A. I draw the hops and prioritise. Functional analyst or Risk SME clarifies expected behaviour. Data engineers (Brussels and Warsaw) build and reconcile. Credit Risk accepts. I keep audit evidence if capital-relevant. Scrum Master facilitates; they do not own priority.

Q6. How do you prioritise when Risk, Data Management and IT all say urgent?

A. Put trade-offs on the backlog: mandatory regulatory, capital/LLP impact, data defect blocking engines, then improvements. One early conversation with the right people from each tribe.

Q7. What does Risk Data Journey mean to you here?

A. How credit risk data moves from need and sources through the Data Lake into PD/EAD/LGD and related engines, then out to capital, RWA, LLP and tooling for Risk and Business, with lineage auditors can follow.

Q8. How would you support a regulatory initiative in the Data Lake?

A. Clarify fields and timing with Credit Risk; align Data Management on landing and quality; sequence engine consumption after lake readiness; keep an audit trail.

Q9. Oracle and SAS Enterprise Guide?

A. Hands-on from early banking roles on risk reporting and credit risk data. Later stacks differed, but Oracle and SAS are foundation I built in regulated risk data. Ready for ING's stack.

Q10. Record Bank data landscape?

A. Residual credit data still matters after integration. I do not claim day-one Record table knowledge. I ramp by mapping sources, joins and engine impact with Risk and Data Management.

Q11. How do you work with auditors?

A. JD asks for key contact. Prepare lineage, reconciliations and clear status with Risk SME and engineer. Backlog control gaps instead of defending them.

Q12. Your title says Director / CoE. Will you do hands-on PO work?

A. Yes. I want long-term ownership on one credit risk engine domain: backlog, squad, landscape, auditors.

Q13. What would you digitalize first?

A. Whatever still forces Risk or Business into manual reconciliations around engine outputs. Thin automated flow with clear acceptance, then expand.

Q14. How do you measure squad success here?

A. Regulatory windows met, engine stability, clear backlog across tribes, fewer defects escaping to LLP/RWA noise, audit questions answered without fire drill.

Q15. Why ING Credit Infrastructure?

A. It sits where I am strongest: credit risk engines, PD/EAD/LGD landscape, Risk, Data and IT bridge, and PO ownership under regulation. KBC is the closest twin. Belgium is where I am returning.

14. Day-in-the-life Q&A (on the job)

Q16. Stand-up: engine batch failed overnight. What do you do?

A. Confirm impact with engineer. If LLP/RWA path is at risk, notify Credit Risk and Finance early. Decide hotfix vs planned story. Update the board.

Q17. Risk wants a change mid-sprint. Data Management says no capacity.

A. Make the constraint visible. If mandatory, escalate jointly and renegotiate sprint scope. If not, park with a date. I do not promise both without capacity.

Q18. Engineer says “data, not engine.” Risk says “engine is wrong.”

A. Journey expert hat: follow lineage (source → lake → engine input → calculation). Assign the fix to the right tribe/squad. Update acceptance so blame stops.

Q19. How do you write a good engine story?

A. Business/regulatory why; data inputs; engine behaviour change; acceptance including reconciliation; lake/Data Management dependency if any; audit evidence note if capital-relevant.

Q20. SteerCo asks for status in two sentences.

A. “Sprint focus is lake field X plus DoD/PD consumption for deadline Y. Risk: Data Management dependency; mitigation: joint refinement tomorrow and cut scope if not ready by Wednesday.”

Q21. Chapter meeting vs squad work?

A. Squad: delivery and backlog for Credit Infrastructure. Chapter: craft standards for POs/journey experts. Capital deadlines stay in the squad/tribe path.

Q22. Modelers want ten mart features. Squad has three slots.

A. Ask which feature unblocks monitoring or next validation. Take one or two. Keep engines stable if a policy deadline is live.

Q23. Business tribe asks for a dashboard you do not own.

A. If it is thin tooling on engine outputs in JD scope, shape a small story. If it is portfolio reporting owned elsewhere, connect them to the reporting squad and keep your backlog clean.

Q24. You are new. How do you learn the tribe map fast?

A. Ask for the org sketch on day one. Meet Credit Risk Area SME, Data Management PO, chapter lead, Warsaw lead. Draw the engine map. Sit in one batch failure and one lake change. Then own a small cross-tribe story.

Q25. Conflict: digitalization roadmap vs this month’s regulatory patch.

A. Patch and compliance first. Keep one digitalization slice alive if capacity allows. Explain that openly to Credit Risk and Business stakeholders.

Q26. Who is “Accountable” when a lake field is late?

A. Data Management owns landing. I own making the dependency and engine impact visible. Credit Risk owns whether the business date can move. We escalate together, not in silence.

Q27. How do Warsaw hubs change your PO style?

A. Stories must be executable without hallway context. Shared glossary, examples, and acceptance. Pair early on complex lineage. Respect ceremony times across locations.

15. One-minute pitch (HM open)

Credit Infrastructure sits in Risk Platforms under Integrated Risk Management and runs the engines that turn credit risk data into parameters for capital, RWA and LLP. Around it sit the Credit Risk Tribe as demand, Data Management as the lake and quality partner, Business tribes as consumers of indicators and tooling, and Audit on evidence. My role is to lead the squad as Product Owner and to master the credit risk data journey across those tribes: PD, EAD, LGD and related engines, Data Lake regulatory work, flows and tooling, and digitalization. That is the work I have done on calculation engines at KBC and in credit risk data at HSBC and Credit Suisse. I am ready to do it hands-on at ING Belgium.

16. Risk keyword glossary (with formulas)

Interview crib sheet. Formulas are standard Basel / IFRS textbook forms. Bank engines implement policy-specific variants. Confirm ING naming (especially DTAI) on day one. Not bank policy

Core model parameters (engines you own)

Term	Definition	Formula / expression
PD Probability of Default	Likelihood that an obligor (or facility, per model design) defaults within a horizon (often 12 months for IRB; lifetime for IFRS 9 Stage 2/3 logic).	$0 \leq PD \leq 1$ Often from score/rating: $PD = f(\text{score, segment, macro})$ Portfolio: average PD weighted by exposure
EAD Exposure at Default	Estimated exposure amount if default occurs. Includes drawn balance plus expected drawdown of undrawn limits.	$EAD \approx \text{Drawn} + CCF \times \text{Undrawn}$ CCF = Credit Conversion Factor (0 to 1) Product-specific caps and netting may apply
LGD Loss Given Default	Share of EAD expected to be lost after default, net of recoveries and collateral, usually after costs and discounting.	$LGD = 1 - \text{Recovery Rate}$ Recovery Rate $\approx \text{Recoveries (collaterals + cures - costs)} / EAD$ $0 \leq LGD \leq 1$ (floors/caps in IRB)
EL Expected Loss	Average loss you expect over the horizon. Feeds LLP thinking; capital covers unexpected loss beyond EL.	$EL = PD \times EAD \times LGD$ (same building blocks your engines produce)
CCF Credit Conversion Factor	Fraction of undrawn commitment expected to be drawn at default.	Used inside EAD: see EAD formula above

Capital, RWA, provisions

Term	Definition	Formula / expression
RWA Risk-Weighted Assets	Exposure scaled by risk so capital ratios compare banks. Higher PD/LGD (and lower collateral quality) usually raise RWA.	IRB idea: $RWA = 12.5 \times K \times EAD$ K = Basel capital requirement from PD, LGD, maturity, correlation $12.5 = 1 / 8\%$ (inverts 8% capital ratio)
K (IRB)	Regulatory capital charge per unit EAD from the Basel IRB formula (unexpected loss style).	$K = f(PD, LGD, M, R) - PD \times LGD$ (simplified: UL capital after EL offset; exact Basel formula is longer)
LLP Loan Loss Provisions	Accounting stock/flow of credit loss allowances on the book (IFRS 9 ECL in modern banks).	Linked to ECL (below). Engines' PD/LGD/EAD (and staging) feed provisioning models
ECL Expected Credit Loss (IFRS 9)	Probability-weighted credit losses. Stage 1: 12-month ECL. Stage 2/3: lifetime ECL.	$ECL \approx \sum (PD_{\text{scenario}} \times EAD_{\text{scenario}} \times LGD_{\text{scenario}} \times \text{discount}) \times \text{weight}$ Stage drives 12m vs lifetime PD
CET1 / Capital ratio	Core capital adequacy vs RWA.	$CET1 \text{ ratio} = CET1 \text{ capital} / RWA$

Rule-based engines and lending indicators

Term	Definition	Formula / expression
DoD Definition of Default	Policy/regulatory rules that mark an obligor or facility as defaulted (e.g. 90 days past due, unlikely-to-pay, specific forbearance cases). Rule engine, not a PD model.	Default flag = 1 if any DoD trigger true Examples: $DPD \geq 90 \cdot UTP$ · certain forbearance (exact triggers = Credit Risk policy)
EWS Early Warning Signal(s)	Pre-default watch indicators from behaviour or financial stress so Risk/Business can act before DoD.	EWS flag / score = f(arrears trend, utilisation, covenants, ...) Thresholds set with Credit Risk
LTV Loan to Value	How large the loan is versus collateral value. Higher LTV → usually weaker collateral cover for LGD.	LTV = Loan amount (or exposure) / Property (collateral) value
DSTI Debt Service to Income	Affordability: debt service payments versus income. Used in origination and monitoring.	DSTI = Debt service (period) / Income (same period)
DTAI	Named in sibling Credit Infrastructure JD with DSTI/LTV. Local name: confirm at ING. Common readings: Debt-to-Assets Indicator, or Debt-to-Available-Income style ratio.	If debt-to-assets: DTAI = Total debt / Total assets If debt-to-available-income: Debt / Available income Ask SME which ING uses
DPD Days Past Due	How long payment is overdue. Key DoD input.	DPD = days since due date unpaid (product rules apply)
UTP Unlikely to Pay	Judgemental/default trigger when repayment is unlikely even if not yet 90 DPD.	Qualitative + quantitative triggers in DoD policy (no single universal formula)

Portfolio, accounting, regulation (you will hear daily)

Term	Definition	Formula / notes
Forbearance	Concession to a distressed borrower (reschedule, temporary lower rate). Can drive DoD/EWS and staging.	Policy flags → engine inputs (no single formula)
NPL Non-Performing Loan	Exposure in default / non-performing per regulatory definition (aligned with DoD in many banks).	NPL ratio $\approx$ NPL stock / Total loans
Cure	Return from default/non-performing to performing after recovery actions.	Affects LGD recovery paths and historical LGD calibration
Collateral coverage	How much exposure is covered by eligible collateral after haircuts.	Coverage $\approx$ Haircutted collateral value / EAD (feeds LGD)
IRB Internal Ratings-Based	Basel approach using bank PD/LGD/EAD models for RWA (vs standardised risk weights).	Uses PD, LGD, EAD, M in Basel K formula
IFRS 9	Accounting standard for ECL staging (1/2/3) and provisions.	See ECL formula; staging $\neq$ DoD but linked
CRR / CRD	EU Capital Requirements Regulation / Directive: legal frame for RWA and own funds.	Policy source for capital rules
AnaCredit	ECB granular credit data reporting. Needs clean facility/counterparty data journeys.	Reporting requirement; engines/data must align
Record Bank residual	Legacy product/codes from Record Bank still present in ING Belgium credit data map.	Mapping problem in landscape (plus on JD)

One-line chain (memorise)

EL = PD × EAD × LGD · EAD $\approx$ Drawn + CCF × Undrawn · LGD = 1 – Recovery · LTV = Loan / Value · DSTI = Debt service / Income · RWA $\approx$ 12.5 × K × EAD · CET1 ratio = CET1 / RWA · DoD/EWS = rule flags that feed models and monitoring

End of simulation. Sources: public JD REQ-10118811 · sibling Credit Risk Data Engineer (engines DoD/EWS/PD/LGD, Data Management, Risk & Business tribes) · public Credit Risk Tribe lead ING Belgium. Use with CV: anjish-bhondwe-cv-ing-po-risk-data-journey-expert.pdf